

richardwu.ca

CS 466/666 COURSE NOTES

DESIGN AND ANALYSIS OF ALGORITHMS

ANNA LUBIW • FALL 2018 • UNIVERSITY OF WATERLOO

Last Revision: September 18, 2018

Table of Contents

1	September 10, 2018	1
1.1	Overview	1
1.2	Travelling Salesman Problem (TSP)	1
1.3	Approach to NP-complete problems	2
1.4	Metric TSP	2
2	September 12, 2018	4
2.1	Data structures	4
2.2	Priority queue	4
2.3	Prim's algorithm	5
2.4	Binomial heaps	5
3	September 17, 2018	7
3.1	Amortization	7
3.2	Amortization "potential" method	7
3.3	Summary of mergeable heaps	9
3.4	Lazy binomial heaps	9

Abstract

These notes are intended as a resource for myself; past, present, or future students of this course, and anyone interested in the material. The goal is to provide an end-to-end resource that covers all material discussed in the course displayed in an organized manner. These notes are my interpretation and transcription of the content covered in lectures. The instructor has not verified or confirmed the accuracy of these notes, and any discrepancies, misunderstandings, typos, etc. as these notes relate to course's content is not the responsibility of the instructor. If you spot any errors or would like to contribute, please contact me directly.

1 September 10, 2018

1.1 Overview

How to design algorithms Assume: greedy, divide-and-conquer, dynamic programming

New: randomization, approximation, online algorithms

For one's basic repertoire, assume knowledge of basic data structures, graph algorithms, string algorithms.

Analyzing algorithms Assume: big Oh, worst case asymptotic analysis

New: amortized analysis, probabilistic analysis, analysis of approximation factors

Lower Bounds Assume: NP-completeness

New: hardness of approximation

1.2 Travelling Salesman Problem (TSP)

Given graph (V, E) with weights on edges $W : E \rightarrow \mathbb{R}^{\geq 0}$ find a *TSP tour* (i.e. a cycle that visits every vertex exactly once and has minimum weight or $\min \sum_{e \in C} w(e)$).

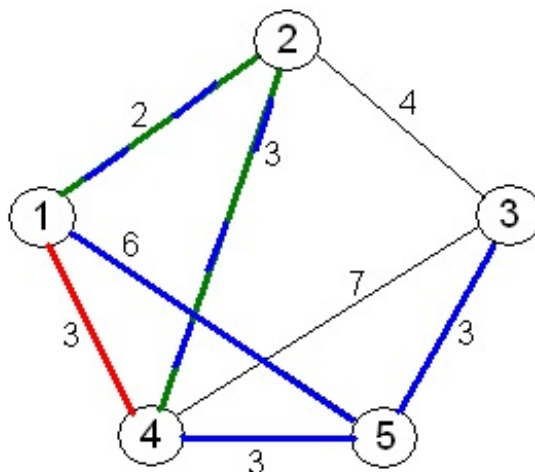


Figure 1.1: TSP tour is highlighted in blue.

Usually assume a **complete** graph (all possible $\binom{n}{2}$ edges exist). We can add missing edges with *high weight* to convert non-complete to complete.

Applications of TSP:

- School bus routes

- Delivery
- Tool path in manufacturing

To show the *decision version* of TSP (exist a tour of total weight $\leq k$) is NP-complete:

1. Show it is in NP (i.e. provide evidence (the tour itself) that there exists a TSP tour and show weights add up to $\leq k$)
2. Show a known NP-complete problem reduces in polynomial time (\leq_p) to TSP (the Hamiltonian cycle problem can be reduced to TSP)

1.3 Approach to NP-complete problems

For NP-complete problems we want to:

- Find exact solutions
- Find fast algorithms
- Solve hard problems

We can in effect only choose two: for hard problems we give up on either *fastness* (exponential time algorithms) or *exactness* (approximation algorithms).

1.4 Metric TSP

An approximation exists for the **metric TSP** version, where:

- $w(u, v) = w(v, u)$
- $w(u, v) \leq w(u, x) + w(x, v) \quad \forall x$

An algorithm (1977) was proposed for metric TSP:

1. Find a minimum spanning tree (MST) of the graph
2. Find a tour by walking *around* the tree.

Think of doubling edges of MST to get Eulerian graph (i.e. every vertex has even degree), which lets us find an Eulerian tour traversing every edge once.



Figure 1.2: Eulerian tour is the solid line around the MST, the dotted lines show shortcuts taken, and the nodes are labelled in order.

3. Take shortcuts to avoid re-visiting vertices

Instead of traversing a node twice (when walking around the MST), we take shortcuts and jump directly to the next unvisited node. By the triangle inequality our path should have a shorter path than if we actually traversed the MST edges twice, i.e.:

$$l \leq 2l_{MST}$$

where l is the length of our tour and l_{MST} is the length of the MST (remember we doubled every edge).

Note the total path length we get will differ depending on which node we start with: thus one must attempt all paths to find the best.

Lemma 1.1. This algorithm is a **2-approximation**, i.e.:

$$l \leq 2l_{TSP}$$

where l_{TSP} is the minimum length of TSP.

Proof. We need to show $l_{MST} \leq l_{TSP}$.

Take the minimum TSP tour. Throw out an edge. This is a spanning tree T . Since

$$l_{MST} \leq l_T \leq l_{TSP}$$

the result follows. □

Exercise 1.1. Show factor 2 can happen.

Analyzing/implementing this algorithm (let n # of vertices, m # of edges):

Steps 2 and 3 take $O(n + m)$.

Step 1 is our bottleneck: we've seen *Kruskal's* (sorted edges with union find to detect cycles) and *Prim's* (add shortest edge to un-visited vertex) MST algorithms.

Prim's took $O(m \log n)$ using a heap. An improvement is using a *Fibonacci heap* (1987) which improves runtime for MST to $O(m + n \log n)$. A further improvement uses a randomized linear time algorithm for finding the MST (1995).

Theorem 1.1. For general TSP (no triangle inequality) if there is a polynomial time algorithm k -approximation for any constant k , then $P = NP$.

Proof. Exercise (hint: start with $k = 2$ and the Hamiltonian cycle problem. Show the 2-approximation can be used to solve the HC problem). \square

Can we improve factor of 2 for metric case? Yes (Christofides 1996):

1. Compute MST
2. Look at vertices of odd degree in MST (there will be an even number). Find a minimum weight *perfect matching* of these vertices.

The MST and perfect matching is Eulerian: take an Eulerian tour and take shortcuts (as before).

Implementation: we need a matching algorithm - the best runtime (in this situation) is $O(n^{2.5}(\log n)^{1.5})$ (1991).

Lemma 1.2. We claim $l \leq 1.5l_{TSP}$. Note that $l \leq l_{MST} + l_M$ (where l_M is the total length of the minimum weight perfect matching). We must show that $l_{MST} \leq l_{TSP}$ and $l_M \leq \frac{1}{2}l_{TSP}$.

Sketch: to show $l_M \leq \frac{1}{2}l_{TSP}$, we show the smallest matching is $\leq \frac{1}{2}l_{TSP}$.

Open question: do better than 1.5 for metric TSP. We know the lower bound is 1.0045 (if we could get 1.0045-approximation then $P = NP$).

There is also the **Euclidean TSP** version where $w(e) = \text{Euclidean length}$. We can get ϵ -approximation $\forall \epsilon > 0$.

2 September 12, 2018

2.1 Data structures

Every algorithm needs data structures. Assume knowledge of:

- Priority queue (heap)
- Dictionary (hashing, balanced binary search trees)

In this course, we look at fancier/better DSES and also amortized analysis.

2.2 Priority queue

Operations supported by a priority queue (PQ) are: insert, delete min (delete), decrease-key, build, merge.

We usually implement PQs with a **heap**: a binary tree of elements where the parent is \leq than the left and right children (min-heap), therefore the min. is at the root.

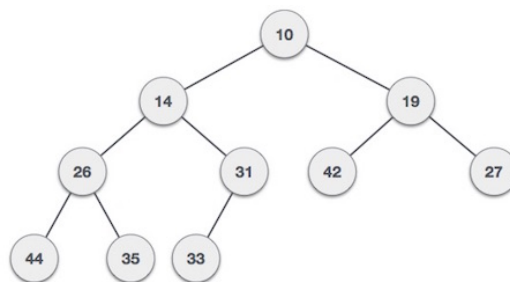


Figure 2.1: An example of a (min-)heap.

We assume that the shape is that of an almost perfect binary tree, where we always add a new element to the bottom right of the tree (if incomplete level) or the bottom left (start of a new level). We can store heaps in *level order in an array*, where for a given element indexed at i , accessing the parent via index $\lfloor \frac{i}{2} \rfloor$ and accessing the children via indexes $2i + 1$ and $2i + 2$.

The height of the tree is obviously $\theta(\log n)$. To implement each operation:

Insert Add new element in last position and bubble/sift up to recover ordering property. $\theta(\log n)$.

Delete min Remove root, it's the minimum. Move last position element to root and bubble/sift down (swap with smaller child). $\theta(\log n)$.

Decrease-key Need only bubble/sift up (if $<$ parent). $\theta(\log n)$.

Build Repeated insertion is $\theta(n \log n)$.

Better approach: from bottom to top (after newly initialized heap in-place) bubble/sift down each element. $\theta(n)$.

2.3 Prim's algorithm

An application of heaps/PQs is for **Prim's MST algorithm**. Given graph with weights on edge, find spanning tree of minimum sum of edge weights.

For our given tree T so far, we find the minimum weight edge connecting to a new vertex. We begin with a PQ of all the edges and we will need to delete any newly added edge and any edges that lead to a newly added vertex.

Let $m = |E|$ number of edges and $n = |V|$ number of vertices. Every edge joins and leaves the heap once for a total of m times. We do delete min n times, so we have $\theta((m + n) \log m)$.

A better approach by using a heap of vertices and keeping track of shortest distance from T to vertex v gives us $\theta((m + n) \log n)$, which is only a constant time improvement since $\log m = \log n^2 = 2 \log n$. We require the decrease-key operation here.

An even better approach for Prim's yields $\theta(m + n \log n)$ via *Fibonacci heaps*.

2.4 Binomial heaps

Binomial heaps improve the merge operation for heaps (which we can use for all other operations).

We use pointers to implement trees and each parent has an arbitrary number of k children (not necessarily binary tree) while maintaining heap order where parent is \leq key of all children. We thus need to relax the shape; we also allow *multiple trees* for a given heap.

We define binomial trees B_k in terms of their rank k , which also coincides with the degree of the root.

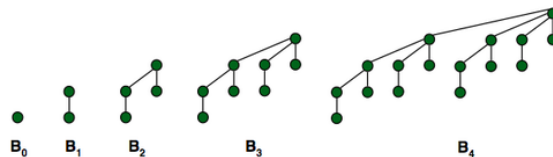


Figure 2.2: Example of five binomial trees with ranks $0, 2, \dots, 4$.

In general, the number of node in B_k is $2^k = 2^{k-1} + 2^{k-1}$.

The height of B_k is k since, by induction, we have the recurrence $height(k) = 1 + height(k - 1)$.

The number of nodes at depth i in B_k is

$$\binom{k-1}{i} + \binom{k-1}{i-1} = \binom{k}{i}$$

(show by induction).

Brief counting proof of binomial equivalence: given k items from which we want to choose i items, we can either pick the first item or not: if we did not pick the first element, then we need to pick i items from the remaining $k-1$ items; if we did pick the first element, we need to pick $i-1$ items from the remaining $k-1$ items.

Note that B_k 's only permit powers of 2 number of elements: thus a **binomial heap** for n elements use a collect of B_i s (heap ordered), at most one for each rank.


E.g. for $n = 13$, we have $13 = 2^3 + 2^2 + 2^0$ (from binary 1101 so we use B_3, B_2, B_0).

It *does not matter* which B_i 's contain a particular element.

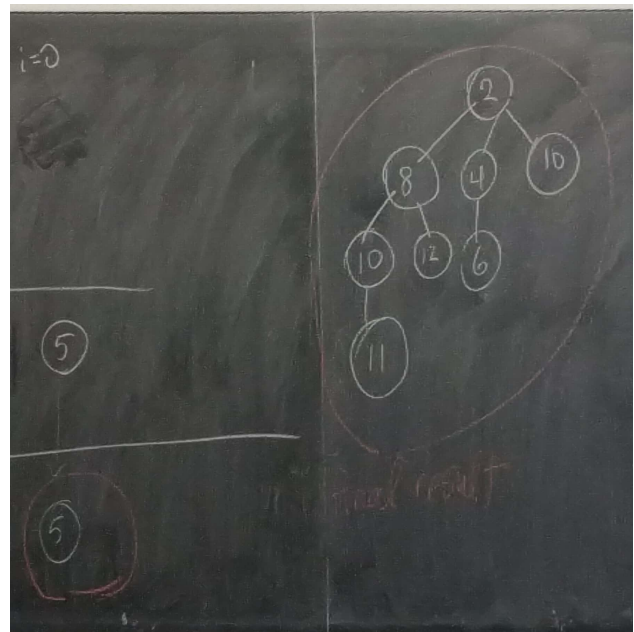
We use $\theta(\log n)$ trees for a given binomial heap with n elements (from binary number expansion).

For merging two binomial heaps, we follow addition of two binary numbers (which are our bitmasks of trees). E.g. for adding a heap with 6 elements to a heap with 3 elements, we add $110 + 11 = 1001$ resulting in a binomial heap with a B_3 and B_0 tree.

When merging two trees of the same rank k , we simply take the tree with the larger root and add it as a children of the root of the other tree.



(a) We carry through the B_0 from the second heap and merge the B_1 trees from two binomial heaps by making the tree with the larger root the children of the other root.



(b) After merging the newly B_2 with the B_2 tree from the first heap into a B_3 tree, we get our final resulting binomial heap with a B_0 and B_3 tree.

Analysis of operations:

Merge Joining two B_i 's take $\theta(1)$. We join up to $\log n$ trees so we have $\theta(\log n)$.

Insert Merge binomial heap with single B_0 (one new element): again $\theta(\log n)$ (since we have the worst case when we insert into a heap with $2^m - 1$ elements).

Delete min Takes $\theta(\log n)$ to find the minimum by checking roots of all trees. Once removing said tree from B_k ,

we end up with $k - 1$ trees (B_{k-1}, \dots, B_0) which we need to merge with the other trees (merge operation is also $\theta(\log n)$) so we have $\theta(\log n)$ overall.

Decrease-key After decreasing key, we bubble/sift up as necessary like before, which takes $\theta(\log n)$ since each tree has height at most $\log n$.

Build Repeated insertion appears to be $O(n \log n)$, but in fact it is $\theta(n)$. This can be seen by repeated addition of 1 in binary to our cumulative total: we only do merges at certain times.

Proof. In general, the cost of incrementing a k -bit counter has a *worst-case cost* of $k + 1$ ($\theta(k)$ bit operations) where you add 1 to $11 \dots 1$.

We show that incrementing from 0 to n is $\theta(n)$:

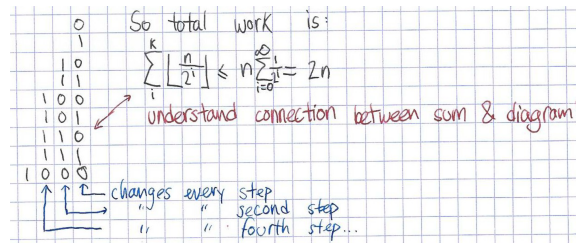


Figure 2.4: A sequence of incrementing from 0 to 1000_2 or 8_{10} . We observe that for binary bit i (rightmost is bit 0) the value of it changes every 2^i step. Thus for incrementing up to n , we sum up the number of times each bit i changes for k bits, which gives us $2n$ changes.

□

3 September 17, 2018

3.1 Amortization

Definition 3.1 (Amortized cost). A sequence of m operations takes total cost $T(m)$: the **amortize cost** of one operation is thus $\frac{T(m)}{m}$.

3.2 Amortization “potential” method

Idea: use an “accounting trick” (“potential” in the physics sense)

Potential “savings” in bank account.

Cost the “true” cost.

Charge artificial: over/under-estimate cost at time of operation.

If charge $>$ cost, we put excess in the bank (add to potential).

If cost $>$ charge, extra has to come out of bank account.

Let Φ_i denote the potential after the i th operation, thus

$$\Phi_i = \Phi_{i-1} + \text{charge}(i) - \text{cost}(i)$$

Theorem 3.1. If the final potential \geq initial potential (almost always 0) then the amortized cost per operation \leq max charge.

Proof. The total charge we've introduced is

$$\begin{aligned}\sum_{i=1}^m \text{charge}(i) &= \sum_{i=1}^m \text{cost}(i) + \sum_{i=1}^m \Phi_i - \sum_{i=1}^m \Phi_{i-1} \\ &= \sum_{i=1}^m \text{cost}(i) + \Phi_m - \Phi_0\end{aligned}$$

Since $\Phi_m - \Phi_0 \geq 0$ then $\sum_{i=1}^m \text{charge}(i) \geq \sum_{i=1}^m \text{cost}(i)$.

Recall that we have for amortized cost

$$\frac{\sum \text{cost}(i)}{m} \leq \frac{\sum \text{charge}(i)}{m} \leq \max \text{charge}$$

□

To do potential analysis, devise potential/charge for each operation such that $\Phi_m \geq \Phi_0$ (the bank is never “in the red”) and max charge is *small*.

Example 3.1. Applying the potential method to the binary counter example, add an extra “\$1” to each basic bit operation to compensate for when we roll over from string of all ones i.e. $\text{charge}(i) = 2$ (1 for the bit operation and “storing” the other 1 for the future).

By the theorem (assuming hypothesis holds), our amortized cost should be 2 per op.

Let's verify, initial potential $\Phi_0 = 0$

counter	cost	charge	potential
0 0 0 0	1		0
0 0 0 1	1	2	1
0 0 1 0	2	2	1
0 0 1 1	1	2	2
0 1 0 0	3	2	1
0 1 0 1	1	2	2

Intuition: potential is equal to the # of ones in counter, so final potential ≥ 0 (initial potential). Note that with cost of $\frac{3}{2}$ this fails.

Claim. Potential = # of ones in binary expansion of counter.

If this claim holds, final potential always ≥ 0 since we have at least one 1 in binary counter.

Proof. Proof by induction. Suppose current counter is

$$\begin{array}{c}01011 \dots 011 \dots 1 \\ \dots 100 \dots 0\end{array}$$

where we have t_i ones at the end.

We thus have

$$\begin{aligned}\phi_i &= \phi_{i-1} + \text{charge}(i) - \text{cost}(i) \\ &= \phi_{i-1} + 2 - (t_i + 1) \\ &= \phi_{i-1} - t_i + 1\end{aligned}$$

Where we zeroed out our t_i ones (subtract) and added one 1. □

While potential method seems harder than previous sum argument, it is much more powerful in general. This gives us $\theta(n)$ since if amortized cost is p (some constant), then n ops cost $pn \in \theta(n)$.

3.3 Summary of mergeable heaps

	binomial heap	lazy binomial heap	Fibonacci heaps
insert	$O(\log n)$	$O(1)$	$O(1)$
delete-min	$O(\log n)$	$O(\log n)$ (amortized)	$O(\log n)$ (amortized)
merge	$O(\log n)$	$O(1)$	$O(1)$
decrease-key	$O(\log n)$ (bubble-up)	$O(\log n)$	$O(1)$ (amortized; improves MST time)
build	$\theta(n)$	$O(n)$	$O(n)$

3.4 Lazy binomial heaps

Idea: be lazy on merge/insert i.e. allow *multiple trees of same size*. Catch up on delete-min: re-combine to form a proper binomial heap (i.e. when delete-min occurs).

For delete-min with lazy binomial heaps:

1. Look at all roots to find min, remove this root.
2. Consolidate trees:

```

1  for rank = 1 to max rank:
2      while there are >= 2 trees of this rank:
3          link them into one tree

```

where max rank is $\theta(\log n)$

Recall that rank = degree of root = height of tree.

It seems the worst case cost is $\theta(n)$ (after inserting n singletons).

Theorem 3.2. Lazy binomial heaps have $O(\log n)$ amortized cost for delete-min and $O(1)$ for insert/merge.

Proof. Let the potential be the # of trees and $\Phi_0 = 0$.

Clearly $\Phi_m \geq 0$ (we never have negative # of trees) so our previous result for the potential method applies.

We need to determine the charge per operation. Recall

$$\text{charge}(i) = \text{cost}(i) + \Phi_i - \Phi_{i-1}$$

Merge cost is 1 (not doing anything), and # of trees is the same (we combine the potentials of the two trees, no additional trees). So we have charge of 1.

Insert cost is 1, # of trees increases by 1, so we have charge of 2.

Delete-min Let r be the degree of the min node ($O(\log n)$), t be the number of trees before delete-min (Φ_{i-1}).

Thus consolidation will be invoked on $t - 1 + r$ number of trees.

Thus the cost will be $\leq t - 1 + r + O(\log n)$, where we have to merge/link at most $t - 1 + r$ times.

$O(\log n)$ is our loop from rank 1 to max rank ($O(\log n)$) and also keeping track of the # of trees of each rank.

Ultimately $\Phi_i \in O(\log n)$ since we end up with a Binomial heap with $O(\log n)$ trees.

We have

$$\begin{aligned}
 \text{amortize cost} & \stackrel{\text{theorem}}{\leq} \max \text{charge} \\
 & \leq \text{cost}(i) + \Phi_i - \Phi_{i-1} \\
 & \leq t - 1 + r + O(\log n) - t \\
 & \leq r + O(\log n) \\
 & \in O(\log n)
 \end{aligned}$$

since $r \in O(\log n)$.

Thus delete-min has $O(n)$ worst case but $O(\log n)$ amortized.

□